

PASSES-HGV: Extending a Fixed-Grid Spectral Framework to 6-DOF Hypersonic Glide and Fractional Orbital Trajectories

Ethan Knox^{*1}

¹Independent Researcher, Champaign, IL, USA

August 4, 2026

Abstract

This paper extends the fixed-grid spectral multi-physics formulation of the companion PASSES framework [5] from a one-dimensional beam to a two-dimensional lifting body, and from a single atmospheric arc to a trajectory spanning boost, exo-atmospheric coast, and hypersonic reentry. Three changes are required. First, the dense Chebyshev collocation operator, whose condition number grows as $\mathcal{O}(N^8)$ for fourth-order bending, is replaced by an ultraspherical spectral discretization [11, 16] in which the differential operators are banded and the conditioning is $\mathcal{O}(1)$; the bivariate extension is assembled as a Kronecker-structured generalized Sylvester system. Second, Euler–Bernoulli kinematics are replaced by a Mindlin–Reissner anisotropic plate formulation [9, 12], which admits the transverse shear deformation that dominates flutter behavior in thicker hull sections, at the cost of requiring three independent free-edge conditions rather than the classical Kirchhoff effective-shear pair. Third, the thermal model is split by material class: charring pyrolysis in the style of CMA [10] for phenolic acreage, and a single-temperature oxidative recession model for non-pyrolyzing refractory leading edges, both mapped to a fixed computational domain by a Landau transformation.

The aerothermal closure couples oblique shock-expansion pressure, Fay–Riddell stagnation convective heating [3] with a modified Lees distribution [6], and Tauber–Sutton radiative heating [15] directly to the deformed structural state. Trajectory generation is posed as a successive convexification problem [7, 14] with free virtual controls under an exact ℓ_1 penalty and a trust region governed by the actual-to-predicted cost reduction ratio. Navigation through plasma blackout uses an ionization-gated Extended Kalman Filter; we give the correct growth rates for the unaided covariance, which are t^3 , t^4 , and t^6 in the velocity-random-walk, accelerometer-bias, and gyro-bias-through-gravity channels respectively, rather than the quadratic growth sometimes assumed. Exo-atmospheric propagation retains the J_2 geopotential term [18].

This manuscript presents the formulation and a falsifiable verification plan. Numerical results are marked in place and are the subject of a companion experimental report.

Keywords: Ultraspherical Spectral Method, Mindlin–Reissner Plate Theory, Hypersonic Aerothermodynamics, Successive Convexification, Plasma Blackout, J_2 Perturbation, Applied Mathematics.

AMS Subject Classification: 65M70, 74K20, 49M37, 70M20.

^{*}Email: ethank5149@gmail.com

Contents

Nomenclature	4
1 Introduction	6
1.1 Scope and relationship to the companion paper	6
2 Related Work	7
3 The Expanded Global State	7
3.1 Attitude parameterization	7
3.2 Local incidence angle on the deformed surface	8
3.3 Aerodynamic forcing on the spectral grid	8
4 Hypersonic Aerothermodynamics	8
4.1 Stagnation-point convective heating	8
4.2 Radiative heating	10
4.3 Heat flux distribution over the planform	10
4.4 Surface recession on a fixed computational domain	11
5 Structural Model: Mindlin–Reissner Anisotropic Plates	11
5.1 The three-field structural state	11
5.2 Constitutive relations and governing equations	12
5.3 Free-edge boundary conditions	12
5.4 Ultraspherical discretization and Kronecker structure	13
6 Guidance and Navigation	14
6.1 Successive convexification	14
6.2 Plasma blackout and ionization-gated estimation	15
6.3 Glide guidance: drag tracking and bank modulation	16
7 Orbital Mechanics and the Coast Phase	16
7.1 J_2 -perturbed propagation	16
7.2 Regime transition within a single integration	17
7.3 Powered hypersonic cruise	17
7.4 Fractional orbital profiles	18
7.5 Midcourse trajectory correction	19
7.6 Dispensing multiple vehicles from a post-boost bus	20
8 Problem Specification in Universal Coordinates	21
9 System Architectures	21
9.1 End-to-end budgets and the closure question	22
10 Verification and Validation Plan	25
11 Limitations	27
12 Conclusion	28
Code and Data Availability	28
Acknowledgments	29

References	31
Appendix A Ultraspherical Operators	32
Appendix B Quaternion Kinematics and Frame Conventions	32

Nomenclature

Symbol	Definition
<i>Geometry and kinematics</i>	
ξ, η	computational coordinates on the planform, $[-1, 1]^2$
$\mathbf{q}$	attitude quaternion, $[q_0, q_1, q_2, q_3]^\top$
$\boldsymbol{\omega}_B$	body angular rate, rad/s
$\mathbf{n}(\xi, \eta, t)$	outward unit normal of the deformed surface
$\hat{\mathbf{v}}_B$	unit freestream velocity in body axes
δ_c	local incidence angle, rad
α	vehicle angle of attack, rad
<i>Structural</i>	
w	transverse deflection, m
ϕ_x, ϕ_y	rotations of the cross-section normal, rad
h	plate thickness, m
$\mathbf{D}_{ij}$	anisotropic flexural rigidity components, N m
G_{xz}, G_{yz}	transverse shear moduli, Pa
κ_s^2	shear correction factor, 5/6
M_x, M_y, M_{xy}	bending and twisting moments, N
Q_x, Q_y	transverse shear forces, N/m
γ_{xz}, γ_{yz}	transverse shear strains
$\mathcal{D}_k$	ultraspherical k -th derivative operator
$\mathcal{S}_k$	ultraspherical basis conversion operator
$\otimes$	Kronecker product
<i>Aerothermodynamics</i>	
q_{dyn}	freestream dynamic pressure, Pa
C_p	pressure coefficient
$\dot{q}_{s,conv}$	stagnation convective heat flux, W/m ²
$\dot{q}_{s,rad}$	stagnation radiative heat flux, W/m ²
Pr, Le	Prandtl and Lewis numbers
β_{Le}	Lewis exponent, 0.52 equilibrium / 0.63 frozen
h_{0e}, h_w	total edge and wall enthalpy, J/kg
h_D	dissociation enthalpy, J/kg
ρ_e, μ_e	boundary-layer edge density and viscosity
ρ_w, μ_w	wall density and viscosity
R_{eff}	effective nose radius, m
$\dot{s}$	surface recession rate, m/s
ΔH_{abl}	effective ablation enthalpy, J/kg
<i>GNC and orbital</i>	
$\mathbf{A}_i, \mathbf{B}_i$	discretized linearized dynamics matrices
$\boldsymbol{\nu}_i$	virtual control vector
w_ν	exact-penalty weight
ρ_k	trust-region actual-to-predicted reduction ratio
$\Delta L^{(k)}$	predicted (convex) cost reduction
n_e	plasma electron number density, m ⁻³
ω_p	plasma angular frequency, rad/s
$\mathbf{P}(t)$	navigation state covariance
J_2	second zonal geopotential harmonic

Symbol	Definition
$\mu, R_{\oplus}$	Earth gravitational parameter and equatorial radius

1 Introduction

Hypersonic glide vehicles and fractional orbital systems present a coupled multi-physics problem that strains conventional computational aerospace methods along several axes simultaneously. The vehicle is a lifting body, so one-dimensional beam kinematics are insufficient. It flies at incidence angles and Mach numbers where aerodynamic loading and aerothermal heating are strongly coupled to the deformed shape, so aerodynamic look-up tables indexed on rigid-body attitude alone lose validity. And it traverses regimes — dense atmosphere, near-vacuum coast, reentry — whose governing balances differ by orders of magnitude, so a formulation tuned to one is typically unusable in another.

The companion paper [5] established a fixed-grid approach to the atmospheric problem: discretize each spatial domain once, hold the grid fixed, and advance the entire coupled system as a single system of ordinary differential equations. That paper’s structural model is a one-dimensional Euler–Bernoulli beam discretized by dense Chebyshev collocation with null-space boundary projection. This paper extends the same principle to a two-dimensional lifting surface and a full-trajectory arc. Three extensions are required, each forced by a specific failure of the one-dimensional formulation.

Conditioning. The dense Chebyshev representation of a fourth-order operator has a condition number growing as $\mathcal{O}(N^8)$. In one dimension at modest N this is tolerable, and [5] manages it by null-space projection. In two dimensions the tensor-product operator inherits the same growth while the problem size grows as N^2 , and the dense operator becomes both ill-conditioned and expensive. This paper adopts the ultraspherical spectral method of Olver and Townsend [11], in which differential operators are banded and the conditioning is $\mathcal{O}(1)$, with the bivariate construction following [16] and the resulting Kronecker-structured system solved by the blocked recursion of [1].

Kinematics. Euler–Bernoulli theory assumes cross-sections remain normal to the deformed axis, excluding transverse shear. For the thicker sections of a waverider hull, transverse shear is not a correction — it is the dominant contributor to the torsional modes that govern flutter. This paper adopts Mindlin–Reissner anisotropic plate theory [9, 12], in which the normal rotations ϕ_x, ϕ_y are independent fields rather than gradients of w .

Regime transitions. A trajectory spanning boost, coast, and reentry passes through conditions where the atmospheric terms vanish and the orbital terms dominate, and back. Section 7 describes how the framework handles this within a single integration.

1.1 Scope and relationship to the companion paper

This manuscript is a formulation paper: it presents the mathematical model, the discretization, and a falsifiable verification plan. Since the plan was written, a reference implementation has been built and the verification tasks executed; where a claim in this paper has been measured, the measured value is quoted inline and attributed, and where a claim proved wrong it has been corrected in place with the correction identified as such (see Remark 4.3). Detailed numerical results, figures, and the full experimental protocol remain the subject of a companion report. Claims not accompanied by a measurement are restricted to what follows analytically or by direct reference to prior work.

Readers should treat [5] as prerequisite for the fixed-grid rationale, the Landau transformation used for ablation, the adaptive filtering architecture, and the Monte Carlo dispersion estimators, none of which are re-derived here. Where this paper’s choices differ from that one, the difference is flagged and justified rather than left for the reader to discover.

2 Related Work

Well-conditioned spectral methods. Chebyshev collocation [17] is the reference discretization for smooth problems on tensor-product domains, but its dense operators condition poorly at high differential order. Olver and Townsend [11] showed that recasting the problem in ultraspherical bases $C^{(\lambda)}$, with sparse conversion operators between them, yields banded differential operators and $\mathcal{O}(1)$ conditioning. Townsend and Olver [16] extended this to bivariate problems on rectangles, expressing a separable-rank- r operator as a generalized Sylvester equation. Chen and Kressner [1] give recursive blocked solvers for exactly this Kronecker structure. This paper’s structural discretization is an application of that line of work to the Mindlin–Reissner system.

Plate theory. Reissner [12] and Mindlin [9] independently introduced first-order shear deformation theories that relax the Kirchhoff normality constraint. The resulting three-field formulation is standard in composite structures; its principal numerical hazard is shear locking as thickness tends to zero, which afflicts low-order finite elements severely. High-order spectral discretizations are substantially less susceptible, and Remark 5.3 states what remains.

Hypersonic aerothermal closure. Fay and Riddell [3] give the stagnation-point convective heating relation for dissociated air including finite-rate diffusion and surface catalysis; Sutton and Graves [13] give a simpler correlation for arbitrary gas mixtures that is widely used where the full Fay–Riddell inputs are unavailable. Lees [6] gives the laminar distribution of heating away from stagnation. Tauber and Sutton [15] give the radiative component, which becomes significant at high entry velocities. Charring ablation follows CMA [10] and FIAT [2].

Convex trajectory optimization. Mao, Szmuk, and Açıkmese [7] established the convergence properties of successive convexification with exact penalty and trust region; Szmuk, Reynolds, and Açıkmese [14] developed the real-time 6-DoF formulation with state-triggered constraints. The formulation in section 6.1 follows these directly.

Positioning. Each component above is established, and no individual model here is novel. The contribution claimed is the coupling: an ultraspherical Mindlin–Reissner structural state, an aerothermal closure evaluated on the deformed surface, and a convex guidance layer, all advanced as one system on grids that never move. Whether that coupling is worth its complexity relative to a conventional partitioned toolchain is an empirical question that section 10 is designed to answer.

3 The Expanded Global State

3.1 Attitude parameterization

Aggressive near-vertical dives and exo-atmospheric tumbling both pass through orientations where three-parameter attitude representations are singular. Attitude is therefore carried as a unit quaternion $\mathbf{q} = [q_0, q_1, q_2, q_3]^\top$ evolving by

$$\dot{\mathbf{q}} = \frac{1}{2} \boldsymbol{\Omega}(\boldsymbol{\omega}_B) \mathbf{q}, \quad \boldsymbol{\Omega}(\boldsymbol{\omega}) = \begin{bmatrix} 0 & -\boldsymbol{\omega}^\top \\ \boldsymbol{\omega} & -[\boldsymbol{\omega} \times] \end{bmatrix}, \quad (3.1)$$

integrated concurrently with the structural state. As in [5], the unit-norm constraint is maintained by Baumgarte stabilization within the right-hand side rather than by post-step renormalization, so that the correction remains visible to the adaptive step-size controller.

3.2 Local incidence angle on the deformed surface

Because the waverider has significant planform curvature and deforms aeroelastically in flight, a single vehicle angle of attack α does not determine the local flow condition at a given surface station. The local incidence angle is defined pointwise from the deformed outward normal $\mathbf{n}(\xi, \eta, t)$ and the unit body-frame velocity $\hat{\mathbf{v}}_B$:

$$\sin \delta_c(\xi, \eta, t) = -\mathbf{n}(\xi, \eta, t) \cdot \hat{\mathbf{v}}_B. \quad (3.2)$$

Remark 3.1 (Sign convention). The negative sign in eq. (3.2) is required and is easy to lose. With $\mathbf{n}$ the *outward* normal, a windward panel — one facing into the flow — has $\mathbf{n} \cdot \hat{\mathbf{v}}_B < 0$. Defining $\sin \delta_c$ without the negation makes windward panels register as leeward, silently inverting the entire pressure and heating distribution. The convention adopted here is $\delta_c > 0$ windward, $\delta_c \leq 0$ leeward.

The normal $\mathbf{n}$ is computed from the deformed surface parameterization, so it depends on the structural state (w, ϕ_x, ϕ_y) ; this is the coupling that makes the aerodynamic load a function of the deformation rather than of rigid-body attitude alone.

3.3 Aerodynamic forcing on the spectral grid

Modified Newtonian pressure alone is inadequate for slender waveriders at moderate incidence: it assigns zero pressure coefficient to all leeward surfaces, discarding leeward suction, and it takes no account of viscous interaction. PASSES-HGV therefore uses a blended closure. On windward panels ($\delta_c > 0$) the local pressure coefficient follows oblique shock-expansion theory, reducing to modified Newtonian

$$C_p = C_{p,\max} \sin^2 \delta_c, \quad C_{p,\max} = \frac{2}{\gamma M_\infty^2} \left(\frac{p_{02}}{p_\infty} - 1 \right), \quad (3.3)$$

where the pressure ratio across the normal shock at the stagnation point is given by the Rayleigh–Pitot relation. On leeward panels ($\delta_c \leq 0$), Prandtl–Meyer expansion is applied until the vacuum limit $C_p = -2/(\gamma M_\infty^2)$ is reached, beyond which the surface is treated as being at vacuum pressure.

The distributed forcing vector on the collocation grid is then

$$\mathbf{Q}_{\text{aero}} = q_{\text{dyn}} \mathbf{C}_p, \quad q_{\text{dyn}} = \frac{1}{2} \rho_\infty V_\infty^2. \quad (3.4)$$

Remark 3.2. Equation (3.3) and its leeward counterpart together are C^0 but not C^1 at $\delta_c = 0$, since the windward branch has zero slope there while the expansion branch does not. This is a genuine violation of the global smoothness the framework otherwise maintains, and it occurs precisely at the shoulder line where panels transition between branches. The two branches are blended over a narrow band $|\delta_c| < \delta_{\text{blend}}$ by a C^2 smoothstep, with δ_{blend} chosen small relative to the incidence variation across a collocation cell. This is a numerical expedient, not physics, and its effect on integrated loads is verification task V4 in section 10.

4 Hypersonic Aerothermodynamics

4.1 Stagnation-point convective heating

The convective heat flux at the stagnation point follows Fay and Riddell [3], which accounts for boundary-layer equilibrium chemistry and the degree of surface catalytic recombination:

$$\dot{q}_{s,\text{conv}} = 0.763 \frac{-0.6}{\text{Pr}} (\rho_e \mu_e)^{0.4} (\rho_w \mu_w)^{0.1} \sqrt{\left(\frac{du_e}{dx} \right)_s} (h_{0e} - h_w) \left[1 + \left(Le^{\beta_{Le}} - 1 \right) \frac{h_D}{h_{0e}} \right]. \quad (4.1)$$

The terms require specification, since eq. (4.1) is frequently quoted without them:

- $\text{Pr} \approx 0.71$ for air; $Le \approx 1.4$.
- $\beta_{Le} = 0.52$ for an equilibrium boundary layer, $\beta_{Le} = 0.63$ for a frozen boundary layer with a fully catalytic wall. The choice matters: the bracketed factor differs by several percent between them, and by considerably more for a non-catalytic wall, for which the bracket reduces toward unity.
- h_D is the enthalpy of dissociation of the free-stream mixture, evaluated at edge conditions.
- The stagnation velocity gradient is supplied by the modified Newtonian estimate

$$\left(\frac{du_e}{dx}\right)_s = \frac{1}{R_{\text{eff}}} \sqrt{\frac{2(p_s - p_\infty)}{\rho_s}}. \quad (4.2)$$

Equation (4.2) makes the coupling to ablation explicit: recession increases R_{eff} , which reduces the velocity gradient, which reduces $\dot{q}_{s,\text{conv}}$ as $R_{\text{eff}}^{-1/2}$. Nose blunting is self-limiting, and capturing that requires the recession model of section 4.4 to feed back into eq. (4.2) within the same time step.

Remark 4.1 (Where the constants in eq. (4.1) actually come from). Implementing eq. (4.1) against Fay and Riddell’s original paper rather than against its restatements produced two findings worth recording.

First, the split exponents are *derivable* rather than conventional. The source’s Eq. (45) writes the flux as $q = (Nu/\sqrt{Re})\sqrt{\rho_w\mu_w}(du_e/dx)_s(h_s - h_w)/\text{Pr}$, and its Eq. (62) correlates $Nu/\sqrt{Re} = 0.67(\rho_s\mu_s/\rho_w\mu_w)^{0.4}\{1 + (Le^{0.52} - 1)h_D/h_s\}$ — a *ratio* across the boundary layer, since the parameter “was found to depend only upon the total variation in $\rho\mu$.” Substituting the second into the first gives the 0.4 from the ratio and a residual $0.5 - 0.4 = 0.1$ from the $\sqrt{\rho_w\mu_w}$, which is exactly the $(\rho_e\mu_e)^{0.4}(\rho_w\mu_w)^{0.1}$ of eq. (4.1). This matters practically: the archived scan of Eq. (63) is too degraded to read the exponents off directly, and the assignment can be recovered instead.

Second, and less comfortably, that same substitution shows every constant in the correlation descending from *one* fitted number carried at two significant figures — the 0.67 of Eq. (58)/(62). The 0.94 printed in Eq. (63) is $0.67/0.71 = 0.9437$ rounded; the $\text{Pr}^{-0.6}$ of the footnote is the correlation’s own $\text{Pr}^{0.4}$ scaling set against the Pr^{-1} of Eq. (45). Backing the eq. (4.1) coefficient out of the paper’s three statements therefore gives three answers: $0.67 \text{Pr}^{-0.4} = 0.7684$ from Eq. (62), $0.94 \text{Pr}^{0.6} = 0.7654$ from Eq. (63), and 0.76 from the footnote as printed — a band of 1.1%. The 0.763 we write above, and which is in general circulation, appears nowhere in the source, though it sits inside that band and is as defensible as any of them.

We are not claiming an error. The spread is around one percent in $\dot{q}_s$, well inside the uncertainty of the edge properties that feed the correlation. We record it because a heat-flux coefficient quoted to three or four significant figures implies a traceability that the source does not support at more than two, and because a reader reproducing our results from the original paper would otherwise find a mismatch and reasonably assume an implementation error. Our implementation names all four constants and defaults to 0.763 for continuity with the literature rather than fidelity to the source.

Remark 4.2. Where the inputs to eq. (4.1) are unavailable — typically in early trajectory screening, before a boundary-layer edge solution exists — the framework falls back to the Sutton–Graves correlation [13], $\dot{q}_s = k_{SG}\sqrt{\rho_\infty/R_{\text{eff}}}V_\infty^3$ with $k_{SG} = 1.7415 \times 10^{-4}$ in SI units for air. This is a correlation, not a theory, and results obtained with it should not be reported as Fay–Riddell results. The distinction is noted because conflating the two is a common error.

4.2 Radiative heating

At the entry velocities relevant to fractional orbital return, shock-layer radiation is not negligible relative to convection. The Tauber–Sutton correlation [15] gives

$$\dot{q}_{s,\text{rad}} = C R_{\text{eff}}^a \rho_{\infty}^b f(V_{\infty}), \quad (4.3)$$

where, for air,

$$C = 4.736 \times 10^4, \quad a = 1.072 \times 10^6 V_{\infty}^{-1.88} \rho_{\infty}^{-0.325}, \quad b = 1.22, \quad (4.4)$$

with a additionally capped by nose radius — $a \leq 0.6$ for $1 \leq R_{\text{eff}} \leq 2$ m and $a \leq 0.5$ for $2 < R_{\text{eff}} \leq 3$ m — and subject throughout to the source’s requirement that $a < 1$. The velocity function f is tabulated and rises steeply, effectively as V_{∞}^8 to V_{∞}^{10} over the correlation’s validity range.

Remark 4.3 (Three properties of eq. (4.3) that are easy to lose). An earlier draft of this paper described the nose-radius exponent as $a \approx 1.0$. That is incorrect: a is a function of both velocity and density, it is capped by nose radius, and the source states explicitly that $a < 1$ must always be met. Over the correlation’s own validity envelope it runs approximately 0.25 to 0.6. The qualitative claim below — that blunting *increases* radiative heating — is unaffected, since $a > 0$; the quantitative exponent is.

Second, eq. (4.3) returns $\dot{q}_r$ in W/cm², not W/m². A framework carrying SI units throughout requires a factor of 10^4 , and omitting it produces a four-order-of-magnitude error in a quantity that still resembles a heat flux.

Third, the tabulated f_E extends from 9 km/s, but eq. (4.3)–eq. (4.4) are stated to apply only for $10 \leq V_{\infty} \leq 16$ km/s, $6.66 \times 10^{-5} \leq \rho_{\infty} \leq 6.31 \times 10^{-4}$ kg/m³, and $0.3 \leq R_{\text{eff}} \leq 3$ m. Interpolating the table below 10 km/s is legitimate; evaluating the correlation there is extrapolation of a fit beyond where its authors validated it, and the implementation refuses it unless asked explicitly.

Note that the R_{eff} dependence in eq. (4.3) is of *opposite sign* to that in eq. (4.2): blunting reduces convective heating but increases radiative heating. The optimal nose radius is therefore a genuine trade rather than a monotone preference, and a framework that models only convection will systematically favor over-blunted geometries. Evaluated on the published correlation at $V_{\infty} = 12$ km/s and $\rho_{\infty} = 3 \times 10^{-4}$ kg/m³, the total stagnation heating has an interior minimum near $R_{\text{eff}} \approx 0.8$ m — inside the source’s own nose-radius envelope, and at a condition where radiation exceeds convection.

4.3 Heat flux distribution over the planform

Stagnation heating is mapped across the bivariate grid using the local incidence angle. For windward nodes a modified Lees distribution [6] governs the local flux:

$$\dot{q}_{\text{total}}(\xi, \eta) = (\dot{q}_{s,\text{conv}} + \dot{q}_{s,\text{rad}}) \sin \delta_c(\xi, \eta) \left(\frac{R_{\text{eff}}}{x(\xi)} \right)^{1/2}, \quad (4.5)$$

where $x(\xi)$ is the running length along the surface streamline from the stagnation point. Equation (4.5) is singular as $x \rightarrow 0$ and is therefore evaluated as $\dot{q}_s$ within the stagnation region $x < R_{\text{eff}}$, matching continuously at $x = R_{\text{eff}}$.

Remark 4.4. Equation (4.5) presumes laminar flow. Transition to turbulence raises local heating by a factor of three to five and is the single largest source of uncertainty in windward heating prediction. The framework carries a transition criterion based on momentum-thickness Reynolds number as a switchable model, and dispersion analyses should sample over transition location rather than fixing it, since the resulting uncertainty typically dominates the aerothermal error budget.

4.4 Surface recession on a fixed computational domain

The companion paper [5] argues against treating ablation as a single-temperature Stefan phase change, on the grounds that charring composites pyrolyze across an extended reaction zone and that a Stefan model omits both the pyrolysis gas enthalpy and the blowing reduction of convective heating. That argument stands and is not revisited here. It does not, however, apply uniformly across an HGV thermal protection system, and this paper therefore splits the model by material class rather than adopting one formulation everywhere.

Acreage TPS (charring). Phenolic-impregnated and carbon-phenolic acreage materials pyrolyze. These regions use the full two-phase charring formulation of [5], comprising Arrhenius decomposition kinetics, in-depth pyrolysis gas transport, and a surface energy balance with the blowing reduction factor $\phi = \ln(1 + 2\lambda B')/(2\lambda B')$. That model is not reproduced here.

Leading edges (non-pyrolyzing). Sharp leading edges and control surfaces on an HGV are typically refractory — carbon-carbon, or an ultra-high-temperature ceramic such as $\text{ZrB}_2\text{--SiC}$. These materials do not pyrolyze. They recede by surface oxidation and mechanical removal, with no in-depth reaction zone and no pyrolysis gas. For these regions a single-temperature moving-boundary energy balance is the physically correct model, not a simplification:

$$\rho_m \Delta H_{\text{abl}} \dot{s}(\xi, \eta, t) = \dot{q}_{\text{total}}(\xi, \eta, t) - \varepsilon \sigma_S T_w^4 - k_{\text{mat}} \left. \frac{\partial T}{\partial n} \right|_{\text{wall}}. \quad (4.6)$$

Remark 4.5 (Resolving the apparent inconsistency). Equation (4.6) is the Stefan-type balance that [5] argues against for charring materials, and its appearance here is deliberate rather than a regression. The two models apply to different material classes. Applying eq. (4.6) to a charring phenolic would understate in-depth heating and overstate surface heat flux by omitting blowing; applying the charring model to a non-pyrolyzing ceramic would introduce decomposition kinetics for a reaction that does not occur. The framework selects per region from the material map.

In both cases the receding domain is mapped to a fixed computational domain by the Landau transformation of [5], so that recession appears as a grid-velocity advection term and no node is created, destroyed, or interpolated. The recession rate updates R_{eff} in eq. (4.2) and eq. (4.3) and the deformed normals $\mathbf{n}(\xi, \eta, t)$ in eq. (3.2), closing the aerothermal-structural loop.

5 Structural Model: Mindlin–Reissner Anisotropic Plates

5.1 The three-field structural state

A hypersonic lifting body is not governed by one-dimensional beam mechanics, and classical Kirchhoff–Love thin-plate theory — which constrains the cross-section normal to remain normal to the deformed midsurface — omits transverse shear deformation. In the thicker hull sections, transverse shear is the dominant mechanism of the torsional modes that set flutter margins. The structural state at each spatial coordinate is therefore the three-field vector

$$\mathbf{u}(\xi, \eta, t) = \begin{bmatrix} w \\ \phi_x \\ \phi_y \end{bmatrix}, \quad (5.1)$$

in which the normal rotations ϕ_x, ϕ_y are independent fields rather than gradients of w . The transverse shear strains are consequently nonzero:

$$\gamma_{xz} = \phi_x + \frac{\partial w}{\partial x}, \quad \gamma_{yz} = \phi_y + \frac{\partial w}{\partial y}. \quad (5.2)$$

Kirchhoff–Love theory is recovered in the limit $\gamma_{xz} = \gamma_{yz} = 0$, i.e. $\phi_x = -\partial w / \partial x$.

5.2 Constitutive relations and governing equations

For an orthotropic laminate with principal axes aligned to the planform coordinates, the moment and shear resultants are

$$\begin{bmatrix} M_x \\ M_y \\ M_{xy} \end{bmatrix} = \begin{bmatrix} \mathbf{D}_{11} & \mathbf{D}_{12} & 0 \\ \mathbf{D}_{12} & \mathbf{D}_{22} & 0 \\ 0 & 0 & \mathbf{D}_{66} \end{bmatrix} \begin{bmatrix} \partial\phi_x/\partial x \\ \partial\phi_y/\partial y \\ \partial\phi_x/\partial y + \partial\phi_y/\partial x \end{bmatrix}, \quad (5.3)$$

$$Q_x = \kappa_s^2 G_{xz} h \gamma_{xz}, \quad Q_y = \kappa_s^2 G_{yz} h \gamma_{yz}, \quad (5.4)$$

with shear correction factor $\kappa_s^2 = 5/6$ following Reissner [12]; Mindlin's $\pi^2/12$ is an alternative differing by under 2%. The rigidity components $\mathbf{D}_{ij}(\xi, \eta)$ form the spatially continuous anisotropic flexural rigidity tensor, constructed by hyperbolic blending across material interfaces so that they remain differentiable.

The governing equations are then

$$\rho h \frac{\partial^2 w}{\partial t^2} = \frac{\partial Q_x}{\partial x} + \frac{\partial Q_y}{\partial y} + q, \quad (5.5)$$

$$\frac{\rho h^3}{12} \frac{\partial^2 \phi_x}{\partial t^2} = \frac{\partial M_x}{\partial x} + \frac{\partial M_{xy}}{\partial y} - Q_x, \quad (5.6)$$

$$\frac{\rho h^3}{12} \frac{\partial^2 \phi_y}{\partial t^2} = \frac{\partial M_{xy}}{\partial x} + \frac{\partial M_y}{\partial y} - Q_y. \quad (5.7)$$

Note that eq. (5.5)–eq. (5.7) are second order in each field, in contrast to the fourth-order Kirchhoff equation. This is the structural reason the three-field formulation is attractive for spectral discretization: the highest derivative order is halved, and with it the conditioning penalty.

5.3 Free-edge boundary conditions

Applying the classical Kirchhoff effective-shear condition to a Mindlin–Reissner state over-constrains the perimeter and induces an artificial boundary layer. Because the three fields are independent, three independent conditions are required per free edge. Along a free edge normal to x :

$$M_x = 0, \quad M_{xy} = 0, \quad Q_x = 0, \quad (5.8)$$

which written out are

$$\mathbf{D}_{11} \frac{\partial \phi_x}{\partial x} + \mathbf{D}_{12} \frac{\partial \phi_y}{\partial y} = 0, \quad \mathbf{D}_{66} \left(\frac{\partial \phi_x}{\partial y} + \frac{\partial \phi_y}{\partial x} \right) = 0, \quad \kappa_s^2 G_{xz} h \left(\phi_x + \frac{\partial w}{\partial x} \right) = 0. \quad (5.9)$$

The third condition is evaluated independently of the geometric slope — this is exactly the freedom the three-field formulation provides, and collapsing it back onto $\partial w/\partial x$ reintroduces the Kirchhoff constraint the formulation exists to relax.

Remark 5.1 (Corner redundancy and how the conditions are imposed). Imposing eq. (5.8) on all four edges produces a rank-deficient constraint set, and the deficiency is specific rather than incidental. The twist condition $M_{xy} = 0$ belongs to the free-edge triple of *both* edges meeting at a corner, so it is imposed twice there; $M_x \neq M_y$ and $Q_x \neq Q_y$ in general, so those are independent. The set is therefore deficient by exactly four, one per corner. A uniformly simply-supported perimeter, used here as a verification configuration because it admits a closed-form Mindlin solution, is deficient by twelve.

Because the deficiency is a property of the boundary configuration rather than of the discretization, the conditions are imposed by null-space projection — the same treatment the companion paper applies to the free-free beam [5] — which degrades gracefully under redundancy where row replacement does not. In implementation the rank is measured rather than assumed, so an inconsistent boundary assembly surfaces as an out-of-range deficiency instead of as quietly wrong frequencies.

Remark 5.2 (Free edges converge algebraically). Free-edge plates carry weak stress singularities at the corners. In consequence the free-free spectrum converges *algebraically* in N , whereas the simply-supported spectrum on the same operator converges exponentially — measured here to 3.7×10^{-10} relative on the first ten modes at $N = 20$ against the closed-form Mindlin solution. This is a property of the boundary-value problem, not of the discretization, and it means a free-free frequency comparison at modest N should not be expected to reach the tolerance an analytic case reaches. Verification of the plate kernel therefore rests on the closed-form simply-supported reference and on manufactured solutions; the free-free comparison is reported alongside without a pass criterion attached.

Remark 5.3 (Shear locking). As $h \rightarrow 0$ the shear stiffness in eq. (5.4) grows as h^{-1} relative to the bending stiffness, and a discretization unable to represent $\gamma_{xz} \rightarrow 0$ exactly develops spurious stiffness — shear locking. This is severe for low-order finite elements and motivates the reduced-integration and mixed formulations standard in that setting. High-order spectral discretizations are markedly less susceptible, because the polynomial space is rich enough to represent the constrained limit closely. "Less susceptible" is not "immune": locking behavior should be measured against thickness ratio rather than assumed absent, which is verification task V2 in section 10. Measured on the present discretization at $N = 16$ per direction, the fundamental free-free frequency parameter falls from 11.71 at $h/L = 0.2$ — genuine shear softening of a thick section — to a plateau of 13.307, and at $h/L = 10^{-3}$ differs from that plateau by +0.018%, against a 1% criterion. No locking is observed across three decades. What *is* paid as the section thins is conditioning: the shear-to-bending stiffness ratio grows as h^{-2} , and the rigid-body modes separate from the elastic spectrum by correspondingly fewer decades, which is what eventually limits the usable thickness range.

5.4 Ultraspherical discretization and Kronecker structure

Dense Chebyshev differentiation matrices condition as $\mathcal{O}(N^{2k})$ for a k -th order operator, giving $\mathcal{O}(N^8)$ for fourth-order bending. This is the same figure quoted in [5], and it is the reason that paper requires null-space projection.

The ultraspherical method [11] avoids the problem at its source. Representing the solution in the Chebyshev basis T_n but the k -th derivative in the ultraspherical basis $C^{(k)}$, the differentiation operator becomes diagonal with a single nonzero band,

$$\mathcal{D}_k = 2^{k-1}(k-1)! \begin{bmatrix} \underbrace{0 \cdots 0}_k & k & & \\ & & k+1 & \\ & & & \ddots \end{bmatrix}, \quad (5.10)$$

and basis conversion between $C^{(\lambda)}$ and $C^{(\lambda+1)}$ is effected by banded operators $\mathcal{S}_\lambda$. Variable coefficients enter as multiplication operators, which are banded whenever the coefficient has a rapidly decaying Chebyshev expansion — which the hyperbolically blended rigidity fields do by construction. The assembled operator is banded with $\mathcal{O}(1)$ condition number, and boundary conditions are appended as dense rows rather than substituted, preserving the bandedness of the interior.

For the bivariate problem on $[-1, 1]^2$, a separable operator of rank r is written

$$\mathcal{L} = \sum_{j=1}^r \mathbf{A}_j \otimes \mathbf{B}_j, \quad (5.11)$$

so that the discretized system is a generalized Sylvester equation in the matrix of bivariate coefficients [16]. The Mindlin–Reissner system eq. (5.5)–eq. (5.7) assembles as a 3×3 block

operator, each block of the form eq. (5.11), and is solved by the recursive blocked algorithm of Chen and Kressner [1], which exploits the Kronecker structure rather than forming the $\mathcal{O}(N^4)$ dense matrix.

Remark 5.4. The $\mathcal{O}(1)$ conditioning claim is a property of the ultraspherical operator itself and is established in [11]. It does not automatically survive the addition of dense boundary rows, variable coefficients with slowly decaying expansions, or the block coupling of eq. (5.5)–eq. (5.7). We assert only what [11] establishes and measure the rest: the conditioning of the assembled Mindlin–Reissner block operator against N is verification task V1.

Measured, the assembled block operator is essentially $\mathcal{O}(1)$ -conditioned: κ moves from 101 to 108 as N goes from 8 to 18 per direction, a fitted log–log slope of 0.09. The block coupling is therefore benign. Two qualifications matter. First, the result is *scaling-dependent*: measured with column equilibration alone the same operator appears to grow as $N^{1.6}$, because a block operator mixes entries carrying different physical units — transverse shear stiffness in N/m against bending rigidity in N·m — and a one-sided scaling leaves that imbalance in the matrix. Two-sided (Ruiz) equilibration removes it. A condition number quoted for a block operator without stating its scaling is not a well-defined quantity. Second, the boundary rows are *not* benign: the projected pencil actually solved grows far faster than the interior, which is precisely the caveat above and localizes where a preconditioner would have to act.

6 Guidance and Navigation

6.1 Successive convexification

Trajectory generation under thermal and structural constraints is posed as a non-convex optimal control problem and solved by successive convexification [7, 14]. At iteration k the nonlinear dynamics are linearized about the current reference trajectory and discretized, giving

$$\mathbf{x}_{i+1} = \mathbf{A}_i \mathbf{x}_i + \mathbf{B}_i \mathbf{u}_i + \mathbf{z}_i + \boldsymbol{\nu}_i, \quad i = 1, \dots, N-1, \quad (6.1)$$

where $\boldsymbol{\nu}_i \in \mathbb{R}^{n_x}$ are virtual controls.

Remark 6.1 (Virtual controls are free, not sign-constrained). The virtual controls in eq. (6.1) are *unconstrained* real vectors. Constraining them to be non-negative — as is sometimes done by analogy with slack variables in inequality constraints — is incorrect here and defeats their purpose. Virtual controls exist to absorb the artificial infeasibility introduced by linearizing non-convex dynamics, and that infeasibility has no preferred sign: the linearized trajectory may need correction in either direction along any state axis. A non-negativity constraint makes the subproblem infeasible in exactly the cases the virtual control was introduced to rescue, and it invalidates the exact-penalty argument of [7], which requires the penalty function to be exact on a neighborhood of zero in all directions.

Because $\boldsymbol{\nu}_i$ is free, it must be penalized to ensure it vanishes at convergence. The exact ℓ_1 penalty

$$J_\nu = w_\nu \sum_{i=1}^{N-1} \|\boldsymbol{\nu}_i\|_1 \quad (6.2)$$

is added to the objective. The ℓ_1 norm is chosen over ℓ_2 specifically because it is an *exact* penalty: for w_ν larger than a finite threshold related to the dual variables of the original problem, the minimizer of the penalized problem coincides with the constrained minimizer, so $\boldsymbol{\nu}_i = \mathbf{0}$ exactly rather than asymptotically. A quadratic penalty drives $\boldsymbol{\nu}_i$ to zero only as $w_\nu \rightarrow \infty$, degrading conditioning.

A trust region bounds the step, accepted or rejected on the ratio of actual to predicted cost reduction:

$$\rho_k = \frac{J(\mathbf{x}^{(k)}) - J(\mathbf{x}^{(k+1)})}{\Delta L^{(k)}}, \quad \Delta L^{(k)} = L(\mathbf{x}^{(k)}) - L(\mathbf{x}^{(k+1)}), \quad (6.3)$$

where J is the true nonlinear cost and L the convex model cost. The step is rejected and the radius contracted when $\rho_k < \rho_0$; accepted with contraction for $\rho_0 \leq \rho_k < \rho_1$; accepted unchanged for $\rho_1 \leq \rho_k < \rho_2$; and accepted with expansion for $\rho_k \geq \rho_2$. Note that $\Delta L^{(k)} > 0$ by construction whenever the convex subproblem makes progress, so eq. (6.3) is well posed; a nonpositive $\Delta L^{(k)}$ indicates subproblem stagnation and terminates the iteration.

6.2 Plasma blackout and ionization-gated estimation

During entry, the shock layer ionizes and the resulting plasma sheath attenuates or reflects radio-frequency signals, severing GNSS updates. The electron number density n_e is estimated from the Saha equation at post-shock conditions, giving the plasma angular frequency

$$\omega_p = \sqrt{\frac{n_e e^2}{m_e \varepsilon_0}}. \quad (6.4)$$

When $\omega_p \geq \omega_{\text{GNSS}}$ the signal is evanescent in the sheath and the measurement is unavailable. The EKF responds by zeroing the corresponding rows of the observation matrix, $\mathbf{H}_k \rightarrow \mathbf{0}$ on the GNSS block, so that the filter propagates on inertial measurements alone and the update step becomes a no-op for those channels.

Proposition 6.2 (Covariance growth during blackout). *For an unaided strapdown inertial solution, the position error covariance during a blackout of duration t grows as*

$$\mathbf{P}_{rr}(t) \sim \underbrace{\frac{1}{3}q_a t^3}_{\text{velocity random walk}} + \underbrace{\frac{1}{4}\sigma_{b_a}^2 t^4}_{\text{accelerometer bias}} + \underbrace{\frac{1}{36}g^2\sigma_{b_g}^2 t^6}_{\text{gyro bias via gravity projection}}, \quad (6.5)$$

where q_a is the accelerometer white-noise PSD, $\sigma_{b_a}^2$ the accelerometer bias variance, and $\sigma_{b_g}^2$ the gyro bias variance.

Proof. Integrating white acceleration noise twice gives position variance $q_a t^3/3$. A constant accelerometer bias b_a produces position error $\frac{1}{2}b_a t^2$, whose variance is $\frac{1}{4}\sigma_{b_a}^2 t^4$. A constant gyro bias b_g produces attitude error $b_g t$, which misprojects the gravity vector by $g b_g t$, which integrates twice to position error $\frac{1}{6}g b_g t^3$, whose variance is $\frac{1}{36}g^2\sigma_{b_g}^2 t^6$. $\square$

Remark 6.3. Proposition 6.2 corrects a claim that is easy to make loosely. The covariance does not grow quadratically during blackout. Quadratic is the growth of the *position error* from an accelerometer bias; the corresponding *covariance* grows as t^4 , and the gyro-bias channel grows as t^6 and dominates for blackouts beyond a few tens of seconds. The distinction is operationally significant: a guidance layer sizing its pull-up trigger on a quadratic model will under-predict the error growth badly at the durations that matter.

The expanding covariance is exposed to the SCvx layer of section 6.1 as a state-dependent cost, so that the optimizer can trade trajectory efficiency against navigation uncertainty — for instance by commanding a pull-up that reduces n_e enough to reacquire GNSS. Whether this trade is exercised usefully, or merely produces chattering at the blackout boundary, is validation task V6.

6.3 Glide guidance: drag tracking and bank modulation

A lifting entry vehicle has one continuous control that matters over the glide, the bank angle σ , and its magnitude and sign act almost independently. The magnitude sets how much of the lift vector is held vertical, which governs the rate of energy depletion and therefore range; the sign sets which way the horizontal component points, which governs crossrange.

Energy rather than time is the natural independent variable, because terminal accuracy is a question of where the vehicle is when it has slowed down, not of where it is at a particular clock reading. The governing relation is exact rather than a modelling choice: with $e = V^2/2 - \mu/r$ and drag deceleration D , energy is depleted at $\dot{e} = -DV$ while range accumulates at $\dot{R} = V \cos \gamma$, so

$$\frac{dR}{de} = -\frac{\cos \gamma}{D}, \quad R_{\text{to go}} = \int_{e_f}^e \frac{\cos \gamma}{D(e')} de'. \quad (6.6)$$

On a shallow glide holding $|\gamma| < 5^\circ$ the cosine costs under 0.4% and eq. (6.6) reduces to the integral of $1/D$. *Commanding a drag profile is therefore commanding a range*, which is why entry guidance in the Shuttle lineage tracks drag rather than altitude or flight-path angle. We verify the integral against its closed form at constant drag to 10^{-12} relative, and against flown trajectories across a sweep of commanded drag.

The drag tracker closes a proportional-derivative loop on $\cos \sigma$, which is the quantity entering the vertical force balance linearly. Its sign convention reads backwards until one notices the mechanism: flying *above* the reference drag means the vehicle is too deep, so the correction must *increase* vertical lift and therefore *reduce* bank magnitude, raising the vehicle into thinner air. Saturation is applied to the bank angle rather than to its cosine, so that a saturated command remains a physically meaningful attitude instead of an out-of-range cosine silently clipped.

Lateral control is a discrete sign flip on a deadband in heading error. A constant deadband is the obvious choice and is the wrong one. Early in the glide the vehicle is fast and a given heading error converts into very large crossrange, while late in the glide there is too little energy left to correct anything and a reversal costs more in drag disturbance than it buys. Scheduling the deadband wide at high speed and tight at low therefore produces few reversals early and tight control late. Measured on a 2100 km glide at $L/D = 2$, holding a single bank sign accumulates 1304 km of crossrange; reversing on the scheduled deadband reduces that to 76 km at the cost of 24 reversals, a factor of seventeen.

Remark 6.4 (What the glide model does and does not carry). The plant is a three-degree-of-freedom point mass over a *non-rotating* spherical Earth with an exponential atmosphere and a fixed lift-to-drag ratio. Earth rotation is a few percent of crossrange over a long glide and its omission is deliberate rather than an oversight: including it changes the plant but not the guidance structure, and leaving it out keeps the verification properties of section 10 attributable to the guidance law rather than to the fidelity of the plant. A vehicle-specific study would substitute the aerodynamic model of section 3.3 for the fixed L/D , at which point the coupling to the deformed-shape incidence of section 3.2 becomes live.

7 Orbital Mechanics and the Coast Phase

7.1 J_2 -perturbed propagation

Over a fractional orbit, a spherical gravity model accumulates unacceptable secular error. The dominant correction is the J_2 zonal harmonic accounting for Earth oblateness [18]. In Earth-Centered Inertial Cartesian coordinates with $r = \|\mathbf{r}\|$,

$$\mathbf{g}(\mathbf{r}) = -\frac{\mu}{r^3}\mathbf{r} + \mathbf{g}_{J_2}(\mathbf{r}), \quad (7.1)$$

$$\mathbf{g}_{J_2}(\mathbf{r}) = -\frac{3J_2\mu R_\oplus^2}{2r^5} \begin{bmatrix} x(1 - 5z^2/r^2) \\ y(1 - 5z^2/r^2) \\ z(3 - 5z^2/r^2) \end{bmatrix}, \quad (7.2)$$

with $J_2 = 1.08263 \times 10^{-3}$, $\mu = 3.986004418 \times 10^{14} \text{ m}^3/\text{s}^2$, and $R_\oplus = 6378137 \text{ m}$. The secular effects are nodal regression and apsidal precession; over a fractional orbit the accumulated cross-range difference relative to a spherical model is of order kilometers, which is large compared with any meaningful terminal accuracy requirement.

Remark 7.1. Higher harmonics, luni-solar third-body terms, and solar radiation pressure are omitted. For a coast of under one orbital period this is defensible — J_2 exceeds the next zonal term by roughly three orders of magnitude — but the omission should be re-examined before the framework is applied to multi-revolution cases. Atmospheric drag during the coast is retained where the density model is non-negligible, since perigee altitude for a fractional orbital profile may be low enough for drag to be significant.

7.2 Regime transition within a single integration

The framework advances boost, coast, and reentry within one integration rather than handing off between phase-specific propagators. Handoffs are a known source of discontinuity: the state must be reinterpolated at the boundary, and the integrator restarts with no history.

The mechanism is that the atmospheric terms are not switched off at the Kármán line but decay smoothly with the density model, so that above roughly 100 km the aerodynamic and aerothermal contributions to the right-hand side become numerically negligible without any branch being taken. The structural state continues to be integrated during coast, where it rings down under whatever damping the model provides.

Remark 7.2 (Cost of the single-integration choice). This choice is not free, and the trade should be stated honestly rather than presented as a pure gain. Continuing to integrate the full structural state through a coast phase in which it carries no meaningful energy means the step-size controller remains subject to the structural stability constraint — the $\mathcal{O}(N^{-4})$ bound of [5] — during a phase whose rigid-body dynamics would otherwise permit steps orders of magnitude larger. A pragmatic mitigation is to freeze the structural block once its modal energy falls below a threshold and reactivate it on atmospheric re-entry, which reintroduces a switch but confines it to a regime where the frozen state is provably quiescent. Which approach is preferable is verification task V5. Measured over a 300 s coast carrying a 100 rad/s structural mode, freezing the block once its modal energy has decayed to 10^{-6} of its initial value costs $15\times$ fewer right-hand-side evaluations and $15\times$ less wall-clock time, with both strategies conserving specific orbital energy to better than 10^{-8} per orbit. The measurement settles the trade in favor of freezing for coasts of this length, and the margin grows with coast duration because the ring-down time is fixed while the coast is not. The switch is defensible only because the freeze condition is checked rather than assumed.

7.3 Powered hypersonic cruise

Every other trajectory in this framework spends energy the vehicle already has. A cruise vehicle adds energy back, which changes range from an accounting problem into a mass-fraction one. For steady level flight the governing result is Bréguet’s, $R = VI_{sp}(L/D) \ln(m_i/m_f)$, and three consequences are worth recording because each is measurable.

The usual gloss on “logarithmic in mass ratio” is backwards. It is commonly said that fuel shows diminishing returns while L/D does not. Over the range of interest the opposite holds. At Mach 8, doubling L/D from 4 to 8 multiplies range by exactly 2.00, while doubling the fuel

fraction from 0.30 to 0.60 multiplies it by 2.57. The reason is that $\ln(1/(1-f))$ is *convex* in fuel fraction — its derivative $1/(1-f)$ grows — so doubling f more than doubles the logarithm. The genuine diminishing return lives in the *mass ratio*: doubling m_i/m_f adds exactly $\ln 2$ whatever it was before. Those are statements about different variables, and conflating them misprices tankage against aerodynamics.

The cruise-climb is a consequence, not an optimisation. Level flight requires lift to equal weight, so as fuel burns the vehicle must either shed lift by reducing incidence — moving off its best L/D — or climb into thinner air and hold the same aerodynamic state. The second is what Bréguet assumes, and in an exponential atmosphere the climb is exactly $H \ln(m_i/m_f)$, *independent of the vehicle*: 3.03 km for a 30% fuel fraction whatever the wing loading, lift coefficient or L/D . That is small enough to fly and large enough that treating cruise as constant-altitude loses range.

Airbreathing wins only where the atmosphere cooperates. A scramjet’s specific impulse falls with Mach number as the captured air’s own kinetic energy grows relative to the heat that can be added, while a rocket’s is constant. With a representative decay model the two cross near Mach 10.5 for the vehicle considered here: below it the airbreather wins by not carrying oxidiser, above it the rocket’s constant impulse is worth more. The crossover Mach is sensitive to the engine model and should be read as a structure rather than as a number.

Remark 7.3 (Comparing a cruiser with a glider). Absolute ranges are close to meaningless across the two, because they buy range in different currencies — one from fuel, the other from initial energy. The comparable quantity is the *break-even fuel fraction*: how much fuel a cruiser must carry to match a given glide. For the vehicle above against a 7700 km glide that fraction is 0.50, which states the price of matching the glide in the only unit that transfers between the two. Cruise range figures here also exclude the acceleration and descent phases at either end, which subtract; a system-level number must add them back.

7.4 Fractional orbital profiles

A fractional orbital profile inserts into a low parking orbit, coasts through part of one revolution, and deorbits onto an entry interface from which the vehicle glides by section 6.3. Three things appear here that a single ballistic arc does not have.

The Earth turns underneath. Every other element of this framework works in an inertial frame, where that is irrelevant. A ground track is not: over one low-Earth revolution of roughly 88 minutes the planet rotates about 22° , so the parking orbit’s track walks steadily westward and repeats only after many revolutions. This is what lets the entry interface be placed almost anywhere by waiting rather than by manoeuvring. We compute latitude geocentrically, which coincides with geodetic on the spherical Earth used here but differs by up to 0.19° — some 21 km on the ground — once an ellipsoid is introduced; that conversion belongs at the point the ellipsoid enters and not earlier.

The approach azimuth is a free variable. A direct ballistic arc between two points lies in the plane containing them and the Earth’s centre, so its arrival heading is fixed by its endpoints. An orbital profile chooses its own plane, and the spherical-triangle relation $\cos i = \sin A \cos \phi$ then fixes the arrival azimuth from the inclination alone — independent of where the trajectory began. At 40° latitude an inclination of 60° arrives on 40.7° ascending or 139.3° descending; a polar orbit arrives due north. The constraint is only that an orbit never reaches latitudes above its inclination.

Perigee depth sets the character of the entry, and it is the dominant design choice.

The deorbit burn itself is cheap: from a 200 km circular orbit, dropping perigee to 50 km costs 45 m/s against an orbital speed of 7784 m/s, under 0.6%. What the burn buys is timing rather than energy. But how deep the perigee is placed governs everything downstream, and the trade is steep:

perigee (km)	Δv (m/s)	transfer arc (deg)	arc (km)	γ_{entry} (deg)
80	35.9	131.4	14 629	−0.393
50	45.0	108.8	12 117	−0.623
0	60.3	89.1	9 920	−0.884
−100	91.4	69.3	7 711	−1.261
−400	188.3	46.2	5 141	−2.042
−1000	401.0	30.5	3 395	−3.192

A negative perigee is virtual — the vehicle never reaches it — and is simply how a steep entry is specified. An order of magnitude in Δv buys roughly a factor of four in transfer arc. The shallow end is not a free choice either: a grazing perigee gives an entry flight-path angle under half a degree and a transfer arc of 14 600 km, which is most of a hemisphere spent descending.

Total ground range is then the parking arc plus the transfer arc plus the glide, and `fobs_profile` closes that accounting exactly. The three legs trade one-for-one: a glide lengthened by 30° shortens the parking arc by precisely 30°, and a profile whose glide and transfer already overshoot the target is reported as not closing rather than returned with a negative parking arc.

Remark 7.4 (Coplanar deorbit is the only sensible option). The burn is taken coplanar throughout. This is not a simplification of convenience: a plane change at orbital speed costs kilometres per second, against tens of metres per second for the entire deorbit. Any azimuth freedom worth having is therefore bought at insertion, by choosing the parking orbit’s inclination, and not later.

7.5 Midcourse trajectory correction

A coast arc that is merely propagated accurately is not yet a trajectory that arrives anywhere. Injection errors, and the unmodeled accelerations the propagation deliberately omits, accumulate into a terminal miss that must be removed by one or more impulsive corrections. The targeting primitive is Lambert’s problem: given the current position, an aimpoint, and the time remaining, find the conic joining them. We solve it in Izzo’s formulation [4], whose independent variable renders the time of flight monotonic and so removes both the bracketing and the case analysis the classical universal-variable method requires near parabolic and half-revolution transfers.

Two things about applying it are worth recording, because both were discovered by measurement rather than anticipated.

The first is that Lambert’s answer is not the maneuver. Lambert is a two-body solve, while the arc actually flown carries J_2 ; over a half-hour coast the resulting terminal miss is kilometres. That is negligible as a fraction of range and far too large as a correction residual. Lambert is therefore used only to seed a Newton iteration on the terminal-position sensitivity $\partial \mathbf{r}_f / \partial \mathbf{v}_0$, evaluated by central differences of the propagator itself so that it is guaranteed consistent with the dynamics being flown. Two passes bring the true miss below a metre. The seed set must also include the vehicle’s current velocity, not merely Lambert’s two branches: late in an arc the vehicle is nearly collinear with its own aimpoint, a geometry in which the transfer plane is barely determined and both Lambert branches can be nonsense — one case in our sweep returned a 127 km/s “correction”.

The second concerns when to burn, and it corrects an intuition we held going in. Correcting a position error δr with time-to-go t_{go} costs $|\Delta \mathbf{v}| \sim |\delta r|/t_{go}$, which we confirm: the product

of maneuver magnitude and time-to-go is constant along the arc to within about ten percent, and equals the uncorrected miss distance. Fuel therefore argues unambiguously for burning early. We had expected execution error to argue the other way, on the reasoning that a large early burn injects a correspondingly large new error. It does inject one — but that error does not penalise burning early, because it is *independent of when the burn occurs*. Execution error scales with the maneuver, the maneuver scales as $1/t_{go}$, and the terminal sensitivity to a velocity change scales as t_{go} ; the two cancel. Isolating the effect by zeroing the navigation covariance, the execution-induced one-sigma miss moves from 97.6 m to 91.6 m across a nineteen-fold range of time-to-go over which the maneuver itself grows seventeen-fold.

What does favour a late burn is knowledge. A correction computed from an estimate can only remove the error the filter has observed; the remainder survives the maneuver untouched. As the navigation covariance contracts along the arc that residual contracts with it, and nothing pushes back — we find the expected miss falling monotonically to the end of the arc even at a pointing error of 100 mrad. The consequence is that midcourse burn timing is not a scalar optimisation with an interior optimum, as it is often presented, but a genuinely two-objective trade between propellant and terminal accuracy. The implementation reports the two separately rather than combining them, since the exchange rate between them is a mission decision and not a property of the dynamics.

7.6 Dispensing multiple vehicles from a post-boost bus

A bus carrying n independently targeted vehicles must place each on a ballistic arc reaching its own aimpoint. Each release is the targeting problem of section 7.5; what only appears once there is more than one release is the interaction between them. Three measured results are worth recording, and two of them contradict the conventional statement of the problem.

Ordering dominates the maneuvers. The assignment of aimpoints to release slots changes total Δv far more than any refinement of the individual maneuvers does. For four aimpoints spread roughly ± 100 km downrange and ± 60 km crossrange about a common impact point, enumerating all 24 orderings gives a spread of 85% between cheapest and dearest — 405 against 748 m/s — and the natural, index-order sequence is not merely suboptimal but is the *worst* of the twenty-four. Since each evaluation requires a full targeting solve per vehicle, exhaustive search is affordable only for small n ; above that we use a 2-opt exchange and report the result as a local optimum rather than as an optimum.

The footprint anisotropy belongs to the timing freedom, not to the geometry. It is usually said that downrange separation is cheap and crossrange separation expensive. Under a *fixed* arrival epoch this is not so. Displacing an aimpoint 50 km from the nominal impact point at a release 600 s into a 2400 s arc costs 35.1 m/s downrange against 40.4 m/s crossrange — a 15% difference, not a qualitative one. The familiar asymmetry appears only when the arrival epoch is allowed to absorb the displacement: permitting arrival to slip by 10 s reduces the downrange cost to 16.1 m/s, a 54% saving, while leaving the crossrange and radial costs *exactly* unchanged with their optima at zero slip. A plane change cannot be bought with time; a downrange displacement can. The distinction matters because a mission that specifies a common arrival epoch has, without necessarily intending to, priced downrange separation as though it were crossrange.

Accumulated error and terminal dispersion are not the same ordering. Every bus maneuver is imperfect and the next vehicle inherits it, so the bus covariance grows monotonically through the sequence — each maneuver contributes a positive-semidefinite block and none is ever removed. The natural conclusion, that the last vehicle released is the least accurate, is false. A later release also has a shorter remaining flight, and terminal miss scales with that flight

time through $\partial \mathbf{r}_f / \partial \mathbf{v}_0$, so the inherited error has less opportunity to grow into a miss. Over a three-vehicle sequence we measure one-sigma dispersions of 1482, 1983 and 1744 m: rising, then falling. Which vehicle needs the accuracy budget is a result of the schedule and cannot be read off the release order.

8 Problem Specification in Universal Coordinates

Every trajectory module here works in an inertial Cartesian frame, which is right for the dynamics and useless for stating a problem. A launch site and a set of aimpoints are naturally given as latitude, longitude and altitude on a reference ellipsoid, so the framework takes WGS-84 geodetic coordinates as its universal input format and converts at the boundary. Three distinctions are easy to elide and expensive to lose.

Geodetic latitude is not geocentric latitude. The first is the angle of the local normal to the ellipsoid — what a survey or a GPS receiver reports — and the second is the angle of the radius from the centre. They differ by up to 0.19° near 45° , which is 21.4 km on the ground, comfortably larger than any terminal accuracy worth discussing.

The Earth-fixed frame is not the inertial frame. A ground point is fixed in ECEF and moves in ECI, so converting a target to an inertial aimpoint requires the epoch *at which the vehicle arrives*, not the epoch of the calculation. An equatorial site travels 465 m/s through the inertial frame; a one-minute epoch error displaces the aimpoint by 27.9 km. Our aimpoint constructor therefore takes the arrival time and threads it through both the frame rotation and the targeting solve.

The ellipsoid is not a sphere. The polar radius is 21.4 km smaller than the equatorial. The glide and fractional-orbital modules assume a spherical Earth for the guidance structure rather than for the geometry, and a system-level answer converts through this section at the boundary rather than carrying the approximation into a terminal error budget.

Coordinate inversion uses Bowring’s closed form refined twice. Bowring alone is sub-millimetre near the surface but degrades with altitude — we measure 0.32 m at geostationary radius — so the refinement is what makes the accuracy claim hold across the whole range rather than only near the ground. Measured round-trip error afterwards is under 10 nm within the atmosphere and under a micrometre to geostationary. Great-circle range is provided for budgeting and labelled as such: haversine on a sphere of the mean radius errs by up to about 0.5% against the true ellipsoidal geodesic, some 30 km over a 6000 km arc, which is negligible against a glide’s own range uncertainty and far too large to use as a terminal aimpoint.

9 System Architectures

The preceding sections build phases: boost, parking orbit, deorbit, midcourse correction, bus dispensing, glide, powered cruise, ballistic entry, terminal homing. A *system* is an ordered sequence of them, and not every sequence is a system. We state the composition rules explicitly and enumerate from them rather than tabulating the familiar architectures, because a table conflates two very different reasons for an architecture’s absence — nobody builds it, and it is not physically admissible — and those look identical once written down.

The rules concern physics and information rather than doctrine. Deorbit presupposes an orbit, since there is no perigee to lower without one, and symmetrically a parking orbit must be left or there is no entry. Dispensing and midcourse correction are exoatmospheric: a bus separating vehicles inside the atmosphere puts them in each other’s hypersonic wake, and section 7.5 solves a Keplerian boundary-value problem that ceases to describe the plant once aerodynamic forces dominate. Exactly one of glide, cruise or ballistic entry may appear, because each describes the whole atmospheric arc rather than a stage of it; chaining two would place the vehicle in two aerodynamic states at once.

Mixed payloads break one of these rules, and correctly so. A bus can dispense glide vehicles and ballistic reentry vehicles on the same pass, and once it does, terminal regime stops being a property of the *architecture* and becomes a property of each *body*: the two classes fly different atmospheric arcs concurrently. Our first rule set forbade this by requiring exactly one terminal regime, which is right for a uniform payload and wrong in general. The corrected rule admits exactly the pair {glide, ballistic} for a mixed payload and nothing else — a propelled cruiser among them would be a third vehicle class rather than a mix of these two — and reports a request for *the* terminal regime of a mixed architecture as the category error it is, rather than silently returning the first. Because the two arcs are concurrent, their order in the phase list carries no physical meaning; listing glide first by convention avoids enumerating every mixed architecture twice, which would otherwise have inflated the count from 54 to 108.

The payload rule needs stating in both directions. Forward, it is capability: a single body has nothing to dispense, a non-lifting body cannot glide, an unpropelled one cannot cruise. The converse is less obvious and is what keeps the enumeration honest. A glide vehicle *can* fly a purely ballistic entry by declining to use its lift, and a cruiser can glide with the engine off; admitting those enumerates one physical trajectory twice under two payload labels, the second copy carrying a description its own flight contradicts. Requiring an architecture to use the capability that names it removes 52 such duplicates from the enumeration, taking it from 226 sequences to 154.

Payload	Terminal regime	Admissible sequences
single reentry vehicle	ballistic entry	12
multiple reentry vehicles	ballistic entry	36
glide vehicle	glide	12
multiple glide vehicles	glide	36
mixed glide and ballistic	glide + ballistic entry	54
powered cruise vehicle	cruise	4
		154

Eleven of these recur often enough to name: ballistic with a single body and with multiple; boost-glide with one glider and with several; fractional-orbital with each of those four payload classes; mixed glide-and-ballistic in both suborbital and fractional-orbital form; and powered cruise. Each is validated against the rules at import, so an inadmissible reference case fails loudly rather than propagating. The enumeration is closed under the rules — everything it produces independently validates — which is what makes it a derivation rather than a curated list kept in parallel with the rules it is supposed to satisfy.

Remark 9.1 (What the count does and does not mean). One hundred admissible sequences is a statement about *composability* and nothing else. It says the phases chain without contradiction; it says nothing about whether any of them is desirable, affordable, or has ever been built. Most of the multiplicity comes from the ordering freedom among independent exoatmospheric phases — whether a correction precedes or follows dispensing, for instance — which is a real distinction with real consequences for dispersion, and not padding.

9.1 End-to-end budgets and the closure question

Admissibility says a phase sequence composes; it says nothing about whether that sequence reaches a given target from a given launch site. Charging each phase its range, Δv and time — and testing the sum against the geometry of section 8 — answers that, and the accounting has one structural feature worth stating.

Most legs contribute a range that is *computed*: a deorbit transfer sweeps an arc set by perigee depth (section 7.4), a glide covers a distance set by its drag profile (section 6.3), a cruise by its

fuel (section 7.3). What closes an architecture is the one remaining leg that absorbs the difference, and which leg that differs fundamentally between the two families. A fractional-orbital profile absorbs it in the *parking arc*, which costs time and no propellant whatever; the closure test is only whether the remainder is non-negative and fits inside one revolution. A suborbital profile has no such leg, so the remainder falls on *boost*, which is bought with propellant, and the test is against a stated booster capability. Treating these two cases alike is the most common way to get a range budget wrong, so the budget names which leg took the slack.

The second point is that an architecture can fail in two ways that are not opposite ends of one scale. For a launch site at 45°N 60°E and an aimpoint at 38°N 100°W — a required range of 10 571 km — a purely ballistic profile falls 8 271 km short of a 2 000 km booster and is infeasible for it. But a fractional-orbital profile carrying a 6 000 km glide *overshoots*: its fixed legs already cover 13 135 km, and no parking arc can close a negative remainder. The remedy is a shorter glide or a steeper deorbit, not a larger booster; commanding 3 000 km of glide instead closes the profile exactly. Reporting an unsigned magnitude with a named reason, rather than a signed residual, is what keeps those two diagnoses distinct.

For a mixed payload the glide and ballistic arcs are flown concurrently by different bodies, so only the longer is charged as transport. Adding them would double-count a distance the system covers once.

Alongside range and Δv , every architecture reports terminal accuracy as both a circular error probable and a 95% radius. Converting between those two by the circular constant $R_{95}/\text{CEP} = \sqrt{\ln 20 / \ln 2} \approx 2.0789$ is standard practice and is only valid for an isotropic dispersion, which a terminal error ellipse is not: downrange and crossrange errors arise from different mechanisms. Computing the containment radii from the exact elliptical integral rather than from that factor, we find the ratio *rises* with elongation — 2.079 at unit aspect ratio, 2.339 at 2:1, 2.646 at 3.3:1, approaching 2.906 in the one-dimensional limit, which is exactly the normal-distribution value $1.96/0.6745$. Scaling a CEP by the circular constant therefore *under-states* the 95% radius for any real dispersion, by 1.5 to 5% across the architectures tabulated here and by up to 40% at the limit. The containment probability reduces to a one-dimensional quadrature after the radial integral is done analytically, so the elliptical answers are exact rather than a fitted correction to the circular ones.

Reaching a small CEP is not a guidance problem for most of its range. A kilometre-class inertial error cannot be refined away, because it is an error in knowing where the target is relative to the vehicle rather than in flying the trajectory. Closing it needs a terminal sensor, and then enough control authority to act on what the sensor reports. The instinct is that the seeker’s angular accuracy sets the answer; usually it does not. Acquiring at range R with one-sigma σ_θ gives a lateral uncertainty $\sigma_\theta R$ — 20 m for a 1 mrad seeker at 20 km, already adequate — while nulling a 1200 m handover error in the 3.3 s that 20 km buys at 6 km/s demands $N'\Delta/t_{go}^2 \approx 33$ g of lateral acceleration, which a glide vehicle does not have. Acquisition range enters through t_{go}^2 and seeker accuracy only linearly, so doubling the former is worth more than a tenfold improvement in the latter, which our model measures directly.

Four floors can bind, and the achievable miss is the largest; naming which one is the actionable part, because the remedies do not overlap. *Divert-limited* is fixed by reducing the handover error, acquiring earlier, or buying lateral acceleration. *Seeker-limited* is fixed with a better sensor — and the floor is evaluated at the range where the loop stops responding, about one time constant of closing distance, not at acquisition, since taking it at acquisition range is badly optimistic. *Loop-limited* means too few time constants remain for convergence. *Target-location-limited* means the aimpoint’s own coordinates are not known better than that, and no sensor or manoeuvre beats it because the vehicle would be flying accurately to the wrong place.

For the fractional-orbital glide architecture, a 1200 m handover error and the following sequence of changes gives:

Configuration	CEP (m)	R_{95} (m)	Binding floor
inertial only	1263	2785	—
1 mrad seeker at 20 km, 5 g, TLE 50 m	1958	4070	divert
+ divert to 20 g	1316	2736	divert
+ acquire at 50 km	58.9	122	target location
+ target location to 10 m	11.8	24.5	target location
+ decelerate to 2 km/s, TLE 3 m	3.5	7.3	target location

Two things are worth reading off that table. The first seeker row is *worse* than inertial-only, because the fallback multiplier the inertial case uses is optimistic and the computed divert-limited answer is not; adding a sensor a vehicle cannot exploit does not help. And past the third row every remaining improvement is in *survey and timekeeping* rather than in guidance — a 10 m CEP is a statement about knowing where the target is, and section 8 shows how easily that is lost, since confusing geodetic with geocentric latitude displaces an aimpoint by 21.4 km and a one-minute epoch error by 27.9 km.

Two of the error contributions are now derived, and one of them corrected an assumption by a factor of twenty. Propagating a 200 m / 1 m/s injection covariance through the actual Kepler transfer of section 7.4 gives entry-interface principal sigmas of 915/639/630 m, 1281 m RSS — which validates the figure previously assumed, but now derives it. The transfer sensitivity is itself informative: $|\partial \mathbf{r}_f / \partial \mathbf{v}_0|$ runs 850 s at a -400 km virtual perigee against 3484 s at $+50$ km, so a shallow deorbit amplifies injection velocity error four times more than a steep one. The perigee-depth choice that trades Δv against transfer arc trades error amplification in the same direction, and deorbiting shallow is cheap in propellant and expensive in accuracy.

The glide contribution was not so kind. Forty closed-loop glides with a 900 m entry-interface position error, 1 m/s speed, 0.02° flight-path angle and 3% ballistic-coefficient scatter gave a terminal dispersion of 34 km downrange and 35 km crossrange — a fortyfold amplification of the entry error, against the 800 m assumed. The defect is structural rather than numerical: eq. (6.6) makes a drag profile equivalent to a range only if the vehicle flies it from the state it was expected to be in, and section 6.3’s tracker has no range feedback. It flies its profile faithfully and lands wherever that profile reaches. Closing an outer loop on predicted range-to-go — comparing $\int de/D$ against the distance actually remaining and scaling the reference drag to null the difference, as entry guidance in the Shuttle lineage does — reduces this to 11.3 km and 6.7 km, a factor of four overall and five in crossrange. The residual is downrange-dominated and resists further gain because the glide is terminated by a speed gate rather than by arrival; handing over on a combined range-and-energy condition is the next structural fix and is not made here. The consequence for section 9.1 is direct: with an 11.3 km handover rather than 1.2 km, reaching a 100 m CEP requires seeker acquisition at 150 km rather than 50 km.

Remark 9.2 (The homing model is a scaling, not an adjoint). The loop is represented by a single time constant and a navigation ratio. A design study would run the adjoint miss-distance analysis and carry glint, radome refraction and target motion separately. What this is good for is sizing — identifying which of the four floors binds, and by how much — and it is deliberately not presented as more.

Remark 9.3 (The error budget is parametric, and says so). The per-phase dispersion contributions are stated inputs, not derived quantities — order-of-magnitude figures chosen to exercise the accounting, with the exception of dispensing, which is scaled from the bus execution model of section 7.6. A CEP produced here is those numbers plus arithmetic. The arithmetic is verified; the inputs are assumptions, and a computed-looking figure should not be allowed to imply a computed provenance.

Remark 9.4 (The budget is an objective function, not an optimiser). Every leg is evaluated at the parameters it is given. Sweeping glide range, perigee depth or fuel fraction to find a design that closes with margin is the caller’s problem, and deliberately so: an optimiser here would bury the trade it exists to expose. Dispersion is propagated where a model exists — the bus accumulation of section 7.6 — and reported as undefined where none does, rather than assumed zero.

10 Verification and Validation Plan

This section states what is measured, against what reference, and what constitutes failure. The plan was written before any implementation existed, so it is falsifiable in advance rather than reported selectively afterward. Tasks V1–V7 have since been executed against the reference implementation; V8 is executed for the radiative correlation and its implementation properties, with the convective reference-case comparison outstanding. V9–V14 were added after section 7.5, section 7.6, section 6.3, section 7.4 and section 7.3 were written, and are stated in the same spirit: both are checkable without external reference data, because the reference is an independent numerical path through the same physics rather than a published number. That is a weaker claim than validation against measurement and is labelled as such. Measured values appear in the relevant sections above and in the companion repository’s verification reports.

ID	Target	Method and reference	Failure criterion
V1	Ultraspherical operator	κ of the assembled Mindlin–Reissner block operator versus N ; comparison against dense Chebyshev collocation	κ growing faster than $\mathcal{O}(N)$
V2	Shear locking	Free-free plate frequencies versus thickness ratio h/L across three decades; comparison against thin-plate Kirchhoff limit	spurious stiffening above 1% at $h/L = 10^{-3}$
V3	Plate dynamics	Natural frequencies of an isotropic free-free square plate against published Rayleigh–Ritz values; method of manufactured solutions on eq. (5.5)–eq. (5.7)	relative frequency error $> 10^{-5}$ for the first ten modes
V4	Aerodynamic blending	Integrated normal force and pitching moment versus δ_{blend} ; sensitivity of trim solution	trim incidence shifting by more than 0.1° across the blend-width sweep
V5	Regime transition	Coast-phase step size and wall-clock, single-integration versus frozen-structure; energy conservation over the coast	secular energy drift exceeding 10^{-8} per orbit
V6	Blackout navigation	Covariance growth against eq. (6.5); SCvx pull-up trigger behavior at the blackout boundary	measured growth not matching eq. (6.5) exponents; chattering at the boundary

ID	Target	Method and reference	Failure criterion
V7	SCvx convergence	Iteration count and virtual control norm at convergence versus w_ν ; verification that $\boldsymbol{\nu}_i \rightarrow \mathbf{0}$ exactly	$\ \boldsymbol{\nu}_i\ $ not reaching zero for finite w_ν
V8	Aerothermal	Stagnation heating against Fay–Riddell reference cases; recession against a FIAT [2] comparison	heating disagreement $> 5\%$ on reference conditions
V9	Lambert targeting	Terminal arrival error of the solved transfer, propagated through the independent J_2 integrator of section 7.1, over a randomized envelope of geometries, times of flight and both directions of motion; conic invariance checked at both endpoints	relative arrival error $> 10^{-7}$ on any physically flyable transfer; energy or angular momentum differing between the two endpoints by $> 10^{-9}$ relative
V10	Bus dispensing	Every released vehicle propagated to its own aimpoint; ordering search checked against exhaustive enumeration; accumulated-covariance monotonicity checked against reported dispersion	any released vehicle missing its aimpoint by > 1 m; the ordering search returning a cost above the exhaustive optimum where enumeration is affordable
V11	Glide guidance	Range-energy integral $\int de/D$ against its closed form at constant drag; flown range against the same integral’s prediction over a sweep of commanded drag; terminal crossrange with bank reversals against the uncorrected single-sign drift	range integral differing from the closed form by $> 10^{-9}$ relative; flown range not monotone in commanded drag; reversals failing to reduce terminal crossrange by at least an order of magnitude
V12	Fractional orbital profile	Closed-form Kepler deorbit solve against the independent J_2 integrator of section 7.1 on radius, swept angle, speed and flight-path angle; Δv against vis-viva; three-leg range accounting	any transfer quantity differing from the integrated trajectory by more than its stated tolerance; Δv differing from vis-viva by $> 10^{-12}$ relative; parking, transfer and glide arcs failing to sum to the angle to the target

ID	Target	Method and reference	Failure criterion
V13	Cruise scaling	Bréguet range against its analytic scalings: exact linearity in L/D , exact additivity of $\ln 2$ per doubling of mass ratio, and vehicle-independence of the cruise-climb $H \ln(m_i/m_f)$ across dissimilar wing loading, C_L and L/D	range not linear in L/D to 10^{-12} ; a doubling of mass ratio not adding $\ln 2$ worth of range; cruise-climb differing between vehicles of equal fuel fraction
V14	Mission budget and accuracy	Elliptical containment integral against the Rayleigh closed form at unit aspect ratio and against the normal quantile in the degenerate limit; radius/probability round-trip; range ledger closure and failure diagnosis over every named architecture	containment radius disagreeing with its closed-form limits by $> 10^{-4}$ relative; a radius not containing the probability that defined it to 10^{-8} ; any named architecture failing to produce a ledger, a verdict and a CEP/ R_{95} pair

Status. V1–V7 and V9–V14 executed and passing against the criteria as stated. V8 is partially complete. The radiative correlation is verified against the published source (Remark 4.3) and the implementation properties hold, but the 5% criterion is stated against published Fay–Riddell reference conditions, which require tabulated equilibrium-air properties not yet transcribed. The recession half of V8 — shared with Paper I’s V4 — is now executed against measured arc-jet recession rather than against a code: seven conditions of Milos and Chen [8] spanning 107–1102 W/cm² and 2.3–84.4 kPa, at 4% median absolute error against a 27% experimental scatter. Full numerical detail is deferred to the companion experimental manuscript.

11 Limitations

1. **No results.** This is a formulation paper. Conditioning, locking, convergence, and throughput claims are arguments from structure or citations to prior work, not measurements of this implementation.
2. **Engineering-level aerodynamics.** Shock-expansion and modified Newtonian closures are inviscid and local. They do not represent shock–boundary-layer interaction, separated flow, base flow, or shock–shock interference at control surface junctions — the last being a principal driver of localized heating peaks on real vehicles. Numerical smoothness in the structural and thermal solvers does not compensate for error entering through this closure.
3. **Laminar heating distribution.** Equation (4.5) assumes laminar flow; transition is modeled by a switchable correlation with substantial uncertainty, which typically dominates the aerothermal error budget.
4. **First-order shear deformation.** Mindlin–Reissner theory assumes a constant transverse shear distribution through the thickness, corrected by κ_s^2 . For thick sandwich construction

with soft cores, higher-order or layerwise theories are required.

5. **Linear structural response.** Equation (5.5)–eq. (5.7) are linear. Geometric nonlinearity under large deflection and material nonlinearity under thermal softening are outside scope, as is aerothermoelastic buckling.
6. **Rectangular computational domain.** The tensor-product ultraspherical construction requires a topologically rectangular parameterization of the planform. Genuinely multiply-connected geometries, or planforms with re-entrant corners, require domain decomposition that is not addressed here.
7. **Constant plate rigidities in the reference implementation.** Equation (5.3) admits spatially varying $\mathbf{D}_{ij}(\xi, \eta)$, and the assembly is written to accept them, but the implementation exercised here uses constant coefficients. Genuinely bivariate coefficient fields require a low-rank approximation of each coefficient so that the Kronecker structure of eq. (5.11) survives; separable fields are a direct extension, general ones are not.
8. **Boundary conditioning, not operator conditioning, is the limiter.** The measurements reported in section 5.4 show the banded interior behaving as advertised while the bordered system does not. Any claim about the scalability of this discretization to large N is a claim about the boundary treatment, and that has not been addressed by preconditioning here.
9. **Equilibrium chemistry.** Equation (4.1) and the Saha estimate of eq. (6.4) assume equilibrium. At high altitude and high velocity the shock layer is chemically frozen or in nonequilibrium, and both the heating and the electron density will be misestimated.
10. **Simplified gravity and environment.** As stated in section 7.1, only J_2 is retained.

12 Conclusion

This paper extended the fixed-grid spectral formulation of [5] to two-dimensional lifting-body structures and full-trajectory arcs. The dense Chebyshev operator was replaced by an ultraspherical discretization with banded operators and Kronecker-structured bivariate assembly; Euler–Bernoulli kinematics were replaced by a three-field Mindlin–Reissner formulation with three independent free-edge conditions; and the thermal model was split by material class, with charring pyrolysis for phenolic acreage and single-temperature oxidative recession for non-pyrolyzing refractory leading edges, resolving what would otherwise be an inconsistency with the companion paper.

Several points were corrected relative to a naive treatment. Remark 3.1 fixes a sign convention in the local incidence angle whose loss silently inverts the entire load distribution. Remark 6.1 establishes that SCvx virtual controls must be free rather than sign-constrained, and why the ℓ_1 penalty is exact where a quadratic penalty is not. Proposition 6.2 gives the correct t^3 , t^4 , and t^6 covariance growth rates during plasma blackout in place of an assumed quadratic. Remark 7.2 states the cost of single-integration regime transition rather than presenting it as free.

What remains is measurement. Section 10 states in advance what will be measured and what would falsify the design.

Code and Data Availability

The manuscript sources, bibliographies, and citation audit for this work are available at <https://github.com/ethank5149/Parameterized-AeroStructural-State-Estimation-Sandbox>. The implementation, container specifications, and figure-generating scripts will be released under the

MIT license in the same repository. **[PENDING: archive a tagged release with a Zenodo DOI at submission]**

Acknowledgments

[PENDING: acknowledgments, if any]

References

- [1] Minhong Chen and Daniel Kressner. “Recursive blocked algorithms for linear systems with Kronecker product structure”. In: *Numerical Algorithms* 84.3 (2020), pp. 1199–1216. DOI: 10.1007/s11075-019-00797-5.
- [2] Y.-K. Chen and Frank S. Milos. “Ablation and thermal response program for spacecraft heatshield analysis”. In: *Journal of Spacecraft and Rockets* 36.3 (1999), pp. 475–483. DOI: 10.2514/2.3469.
- [3] James A. Fay and Frederick R. Riddell. “Theory of stagnation point heat transfer in dissociated air”. In: *Journal of the Aeronautical Sciences* 25.2 (1958), pp. 73–85. DOI: 10.2514/8.7517.
- [4] Dario Izzo. “Revisiting Lambert’s problem”. In: *Celestial Mechanics and Dynamical Astronomy* 121.1 (2015), pp. 1–15.
- [5] Ethan Knox. *PASSES: A Unified, GPU-Accelerated Spectral Framework for Coupled Aeroelastic, Thermodynamic, and Stochastic GNC Flight Simulation*. Companion manuscript. TODO: insert arXiv identifier once posted. 2026.
- [6] Lester Lees. “Laminar heat transfer over blunt-nosed bodies at hypersonic flight speeds”. In: *Jet Propulsion* 26.4 (1956), pp. 259–269. DOI: 10.2514/8.6977.
- [7] Yuanqi Mao, Michael Szmuk, and Behçet Açıkmeşe. “Successive convexification of non-convex optimal control problems and its convergence properties”. In: *2016 IEEE 55th Conference on Decision and Control (CDC)*. Las Vegas, NV, 2016, pp. 3636–3641. DOI: 10.1109/CDC.2016.7798816.
- [8] Frank S. Milos and Yih-Kanq Chen. “Ablation and Thermal Response Property Model Validation for Phenolic Impregnated Carbon Ablator”. In: *Journal of Spacecraft and Rockets* 47.5 (2010), pp. 786–805.
- [9] Raymond D. Mindlin. “Influence of rotatory inertia and shear on flexural motions of isotropic, elastic plates”. In: *Journal of Applied Mechanics* 18.1 (1951), pp. 31–38. DOI: 10.1115/1.4010217.
- [10] Carl B. Moyer and Roald A. Rindal. *An Analysis of the Coupled Chemically Reacting Boundary Layer and Charring Ablator. Part II: Finite Difference Solution for the In-Depth Response of Charring Materials Considering Surface Chemical and Energy Balances*. NASA Contractor Report CR-1061. Aerotherm Corporation, Palo Alto, CA. Washington, DC: National Aeronautics and Space Administration, 1968.
- [11] Sheehan Olver and Alex Townsend. “A fast and well-conditioned spectral method”. In: *SIAM Review* 55.3 (2013), pp. 462–489. DOI: 10.1137/120865458.
- [12] Eric Reissner. “The effect of transverse shear deformation on the bending of elastic plates”. In: *Journal of Applied Mechanics* 12.2 (1945), A69–A77. DOI: 10.1115/1.4009435.
- [13] Kenneth Sutton and Jr. Graves Randolph A. *A General Stagnation-Point Convective-Heating Equation for Arbitrary Gas Mixtures*. NASA Technical Report TR R-376. NASA Langley Research Center. Hampton, VA: National Aeronautics and Space Administration, Nov. 1971.
- [14] Michael Szmuk, Taylor P. Reynolds, and Behçet Açıkmeşe. “Successive convexification for real-time six-degree-of-freedom powered descent guidance with state-triggered constraints”. In: *Journal of Guidance, Control, and Dynamics* 43.8 (2020), pp. 1399–1413. DOI: 10.2514/1.G004549.
- [15] Michael E. Tauber and Kenneth Sutton. “Stagnation-point radiative heating relations for Earth and Mars entries”. In: *Journal of Spacecraft and Rockets* 28.1 (1991), pp. 40–42. DOI: 10.2514/3.26206.

- [16] Alex Townsend and Sheehan Olver. “The automatic solution of partial differential equations using a global spectral method”. In: *Journal of Computational Physics* 299 (2015), pp. 106–123. DOI: 10.1016/j.jcp.2015.06.031.
- [17] Lloyd N. Trefethen. *Spectral Methods in MATLAB*. Software, Environments, and Tools. Philadelphia, PA: Society for Industrial and Applied Mathematics, 2000. DOI: 10.1137/1.9780898719598.
- [18] David A. Vallado. *Fundamentals of Astrodynamics and Applications*. 4th ed. Hawthorne, CA: Microcosm Press, 2013. ISBN: 978-1-881883-18-0.

Appendix A Ultraspherical Operators

The ultraspherical (Gegenbauer) polynomials $C_n^{(\lambda)}$ satisfy the derivative relation

$$\frac{d}{dx} C_n^{(\lambda)}(x) = 2\lambda C_{n-1}^{(\lambda+1)}(x), \quad (\text{A.1})$$

with the Chebyshev polynomials recovered as the limit $T_n = \lim_{\lambda \rightarrow 0} \frac{n}{2\lambda} C_n^{(\lambda)}$ for $n \geq 1$. Equation (A.1) is what makes $\mathcal{D}_k$ in eq. (5.10) sparse: differentiating raises the basis index by one and shifts the degree by one, so the operator has exactly one nonzero diagonal, in contrast to the dense collocation derivative.

Conversion between adjacent bases follows from

$$C_n^{(\lambda)} = \frac{\lambda}{\lambda + n} \left(C_n^{(\lambda+1)} - C_{n-2}^{(\lambda+1)} \right), \quad (\text{A.2})$$

giving tridiagonal conversion operators $\mathcal{S}_\lambda$. A k -th order differential operator with variable coefficients is assembled as

$$\mathcal{L} = \sum_{j=0}^k \mathcal{M}_j^{(k)} \mathcal{S}_{k-1} \cdots \mathcal{S}_j \mathcal{D}_j, \quad (\text{A.3})$$

where $\mathcal{M}_j^{(k)}$ is the multiplication operator for the j -th coefficient function, expressed in the $C^{(k)}$ basis. Bandwidth of $\mathcal{M}_j^{(k)}$ equals the truncation length of the coefficient's Chebyshev expansion, so slowly converging material property fields degrade the sparsity — which is the practical reason for constructing the rigidity fields by hyperbolic blending rather than by piecewise assignment.

Appendix B Quaternion Kinematics and Frame Conventions

The attitude quaternion $\mathbf{q}_{E2B}$ maps ECI to body axes, with scalar part first, $\mathbf{q} = [q_0, \mathbf{q}_{1:3}]^\top$, and the Hamilton (not JPL) product convention. The direction cosine matrix is

$$\mathbf{C}_E^B = \left(q_0^2 - \mathbf{q}_{1:3}^\top \mathbf{q}_{1:3} \right) \mathbf{I} + 2 \mathbf{q}_{1:3} \mathbf{q}_{1:3}^\top - 2q_0 [\mathbf{q}_{1:3} \times]. \quad (\text{B.1})$$

Mixing Hamilton and JPL conventions produces a transpose error that is easy to introduce and hard to detect, since both yield valid rotation matrices; the convention is stated explicitly here for that reason.